CS294-1 Programming Assignment 1

Richard Hwang, David Huang

February 13, 2013

1 Our Solution

We followed advice in the assignment statement, splitting the text with the regex "[\\s.,();:!?&\"]+" and indexing each word, except for stopwords such as "the", "and", "or", "a", and "of". Afterwards, we constructed a sparse matrix, with words as rows and documents as columns.

2 Smoothing/Backoff

We used Laplace Smoothing and found that an α value of TODO yielded the highest F_1 measure.

3 Performance

For some reason, indexing the words was intolerably slow. We tried using foreach, for loops, String.split() and StringTokenizer, and for loops and split turned out to be the fastest combination. It still took at least 10 minutes to index all the words, though. Constructing the word-document matrix, however, was very fast: about 7 seconds. Building the model, that is, computing all the log probabilities took 3.4 seconds. Finally, validation took 13.6 seconds. Placing a flip and flop around our entire main method, we recorded 0.021549 GFlops on a Macbook Pro with Intel Core Duo.